

# AMS 598: Retail Inventory Management Using Big Data

## Group 2

## Professor: Dr. Song Wu

# Project Motivation

## Motivation:

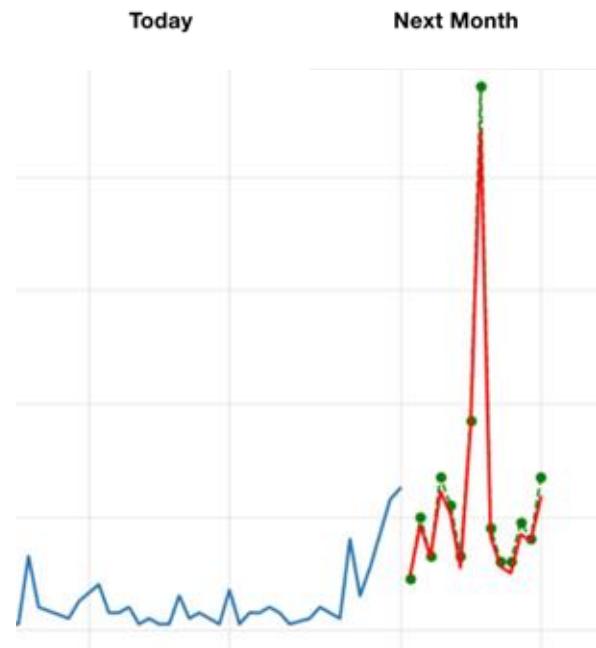
The Consumer Packaged Goods industry relies heavily on data-driven insights to optimize operations and improve decision-making. Accurate forecasting helps minimize lost sales from stockouts and reduce costs from excess inventory.

## Key Decision areas:

- **Portfolio Optimization** – Selecting the right product mix across thousands of stock keeping units.
- **Inventory Management** – Ensuring products are available at the right store, in the right quantity, at the right time.

## Why Big Data?

The datasets contain millions of transactions across many stores and products. Their volume, variety, and velocity require scalable methods to extract insights and support real-time forecasting.



# Problem Statement

## Objective:

Forecast unit\_sales for the next 14 days at the store-item level, using large-scale retail sales data.

## Business Problem:

- Retailers must make daily inventory decisions across thousands of stock keeping units and dozens of stores.
- Stockouts lead to lost revenue, while overstock increases holding costs and waste.
- Accurate short-term demand forecasting is essential for balancing these trade-offs.

## Technical Challenge:

- The dataset contains over 23 million transactions and exhibits strong seasonality, promotions effects, and store item variability.
- Traditional single-machine methods are insufficient for processing and modeling this scale of data.

## Our Goal:

Build a big-data-enabled forecasting system that learns demand patterns, predicts future unit\_sales, and supports inventory optimization.



# Tech Stack and Architecture



## Code Framework:

- GCP, PySpark
- Random Forest forecasting model is implemented using Spark's distributed ML ecosystem.
- All datasets are stored in Parquet format.

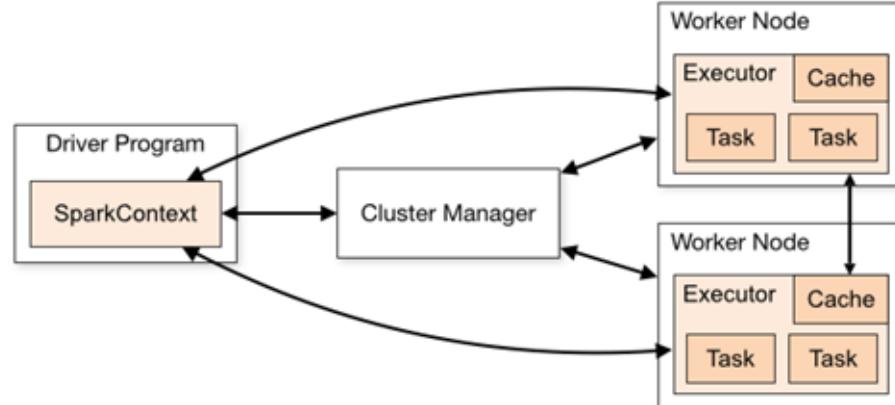
## Phase 1 - Data Wrangling & Preprocessing:

- Distributed Data Ingestion
- Data Cleaning and Sanitization
- Time-Series Feature Engineering
- Dataset Partitioning

## Phase 2 - Model Training, Prediction & Evaluation:

- Distributed Training Pipeline
- Model Persistence & Inference
- Feature Importance Logging
- Prediction
- Performance Evaluation

```
# Initialize Spark with Legacy Parquet Nanoseconds Config
spark = SparkSession.builder \
    .appName(CONFIG["APP_NAME"]) \
    .config("spark.sql.legacy.parquet.nanosAsLong", "true") \
    .getOrCreate()
```



# **Phase 1**

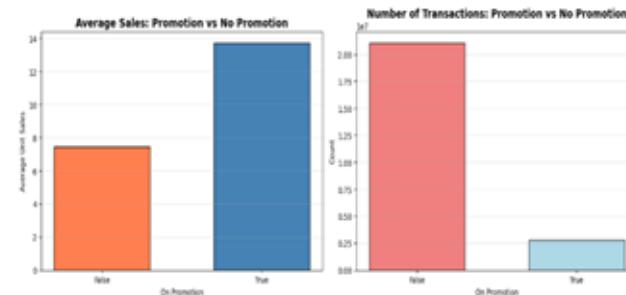
## **Data Wrangling and Preprocessing**

# Dataset Description

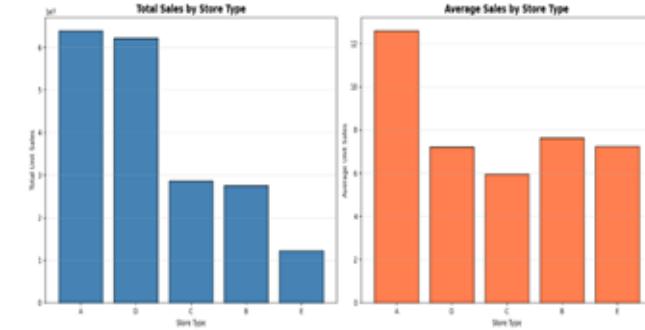
- **Source:** [Corporación Favorita grocery sales dataset](#)
- **Selected:** 2017 subset from the original dataset
- **Format:** Parquet.
- **Scale:** Approximately **23.8 million** daily store–item transaction rows.
- **Scope:**
  - Covers **54 stores** across Ecuador.
  - Includes **4,100 unique products**.
  - **Unit of Data:** Each row represents the sales of a specific product in a specific store on a given day.
- **Time Range:** January 1, 2017 – August 15, 2017.



Promotion Impact on Sales



Store Type Heterogeneity



# Infrastructure: Cluster Configuration - Phase 1

## 1. Architecture: Master-Worker Topology (YARN)

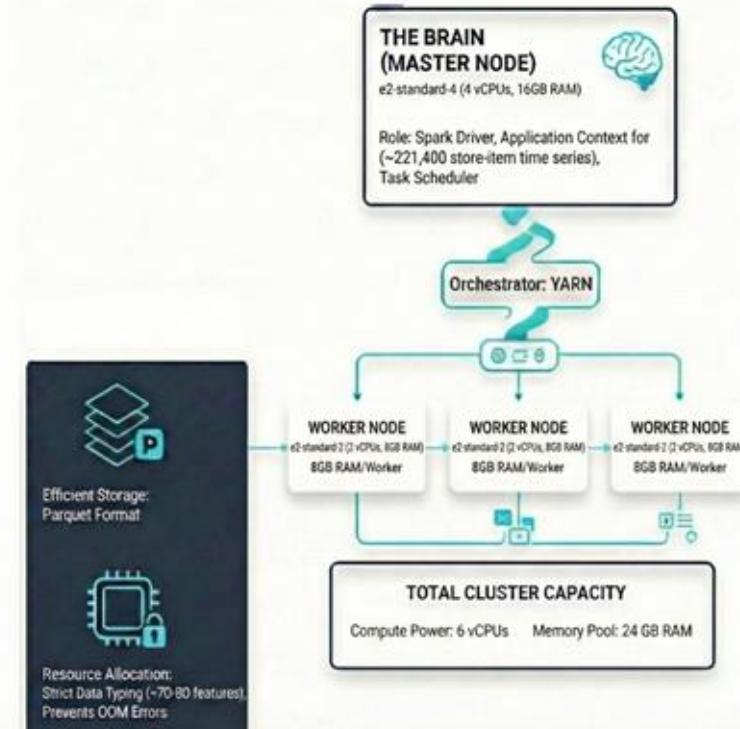
- **Orchestrator:** YARN manages resource allocation across the 3 worker nodes.
- **The Brain (Master Node):**
  - Specs: e2-standard-4 (4 vCPUs, 16GB RAM).
  - Role: Runs the Spark Driver. It maintains the application context for the **~221,400 store-item time series** and schedules tasks (does *not* process heavy data).

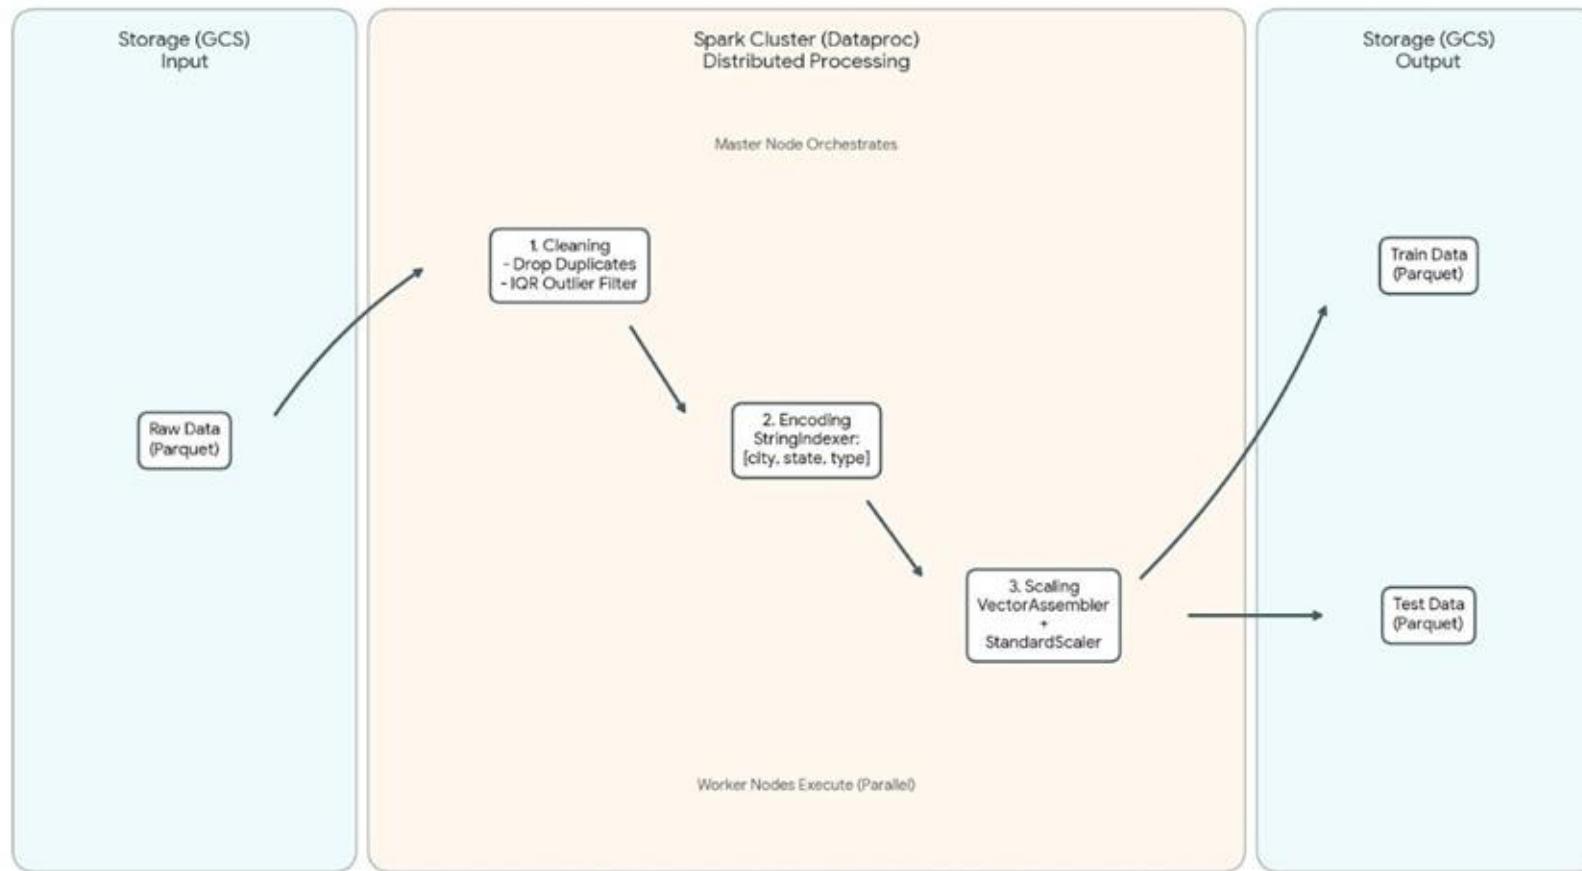
## 2. Total Cluster Capacity (3 Worker Nodes)

- **Compute Power:** 6 vCPUs
- **Memory Pool:** 24 GB RAM

## 3. Memory Optimization for Phase 1

- **Efficient Storage:** Utilization of **Parquet** format allows the smaller cluster to handle the 23.8M row dataset by minimizing I/O overhead.
- **Resource Allocation:** With 8GB RAM per worker, strict data typing was critical to prevent Out-Of-Memory (OOM) errors during the generation of **>70-80 features**.





# Data Preprocessing & Cleaning

## Data Validation:

- Verified schema and data types of all fields using Spark `printSchema()`.
- Checked for missing values across all 20 columns and no missing values found.
- Checked the dataset for duplicate rows using `dropDuplicates()` and no duplicates were detected.

## Data Corrections:

- Replaced negative values in numerical fields (unit\_sales, unit\_sales\_raw, transactions) with 0, since negative values are not meaningful in retail sales.

## Outlier Analysis:

- Performed IQR-based outlier detection on all numerical fields using `approxQuantile()`
- Identified high-outlier columns such as: unit\_sales, unit\_sales\_raw ( $\approx 9.7\%$  outliers), transactions, is\_holiday, class, log\_sales
- Outliers were not removed because large spikes correspond to real retail demand events (e.g., holidays, promotions, payday cycles).

## Categorical Processing:

Used StringIndexer to encode: city, state, type, family.

## Numerical Scaling:

- Created standardized features (Z-score).
- Created min–max normalized features (0–1).

# Feature Engineering

Feature Categories:

Category	Features	Description
<b>Lag Features</b> (7)	t-1, t-2, t-3, t-7, t-14, t-21, t-28	Historical sales at various time lags
<b>Rolling Statistics</b> (12)	Mean, std, min, max for 7/14/30-day windows	Aggregated metrics excluding current row to prevent leakage
<b>EWMA</b> (3)	Spans: 7, 14, 30 days	Exponentially weighted moving averages for trend detection
<b>Change Features</b> (4)	Day-over-day, week-over-week (absolute & %)	Sales differences and percentage changes
<b>Temporal</b> (11)	Year, month, day, DOW, DOY, WOY, quarter, is_weekend, is_month_start/end	Calendar-based features and cyclical patterns
<b>Store-Item Patterns</b> (5)	DOW/month averages, promo lags (7/14 days), 30-day promo frequency	Store-item specific behavioral patterns
<b>Trend</b> (2)	days_since_first_sale, sales_momentum_7	Long-term trends and momentum indicators

**Train-Test Split:** Train: First 213 days | Test: Last 14 days (forecasting horizon) | Target: unit\_sales

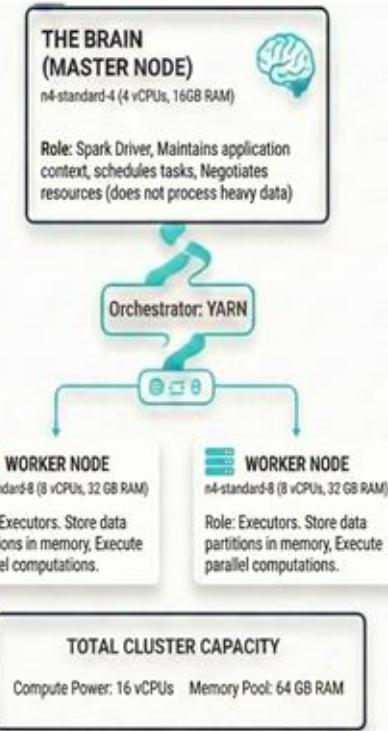
# **Phase 2**

## **Model Training, Prediction & Evaluation**

# Infrastructure: Cluster Configuration - Phase 2

## 1. Architecture: Master-Worker Topology (YARN)

- **Orchestrator:** YARN manages resource allocation.
- **The Brain (Master Node):**
  - **Specs:** n4-standard-4 (4 vCPUs, 16GB RAM).
  - **Role:** Runs the Spark Driver. Maintains application context, schedules tasks, and negotiates resources
- **The Muscle (2x Worker Nodes):**
  - **Specs:** n4-standard-8 (8 vCPUs, 32 GB RAM per node).
  - **Role:** **Executors.** These nodes store data partitions in memory and execute parallel computations.

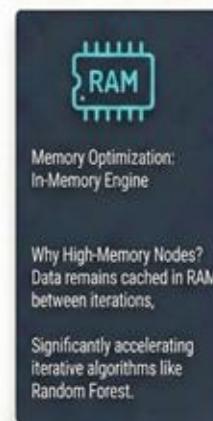


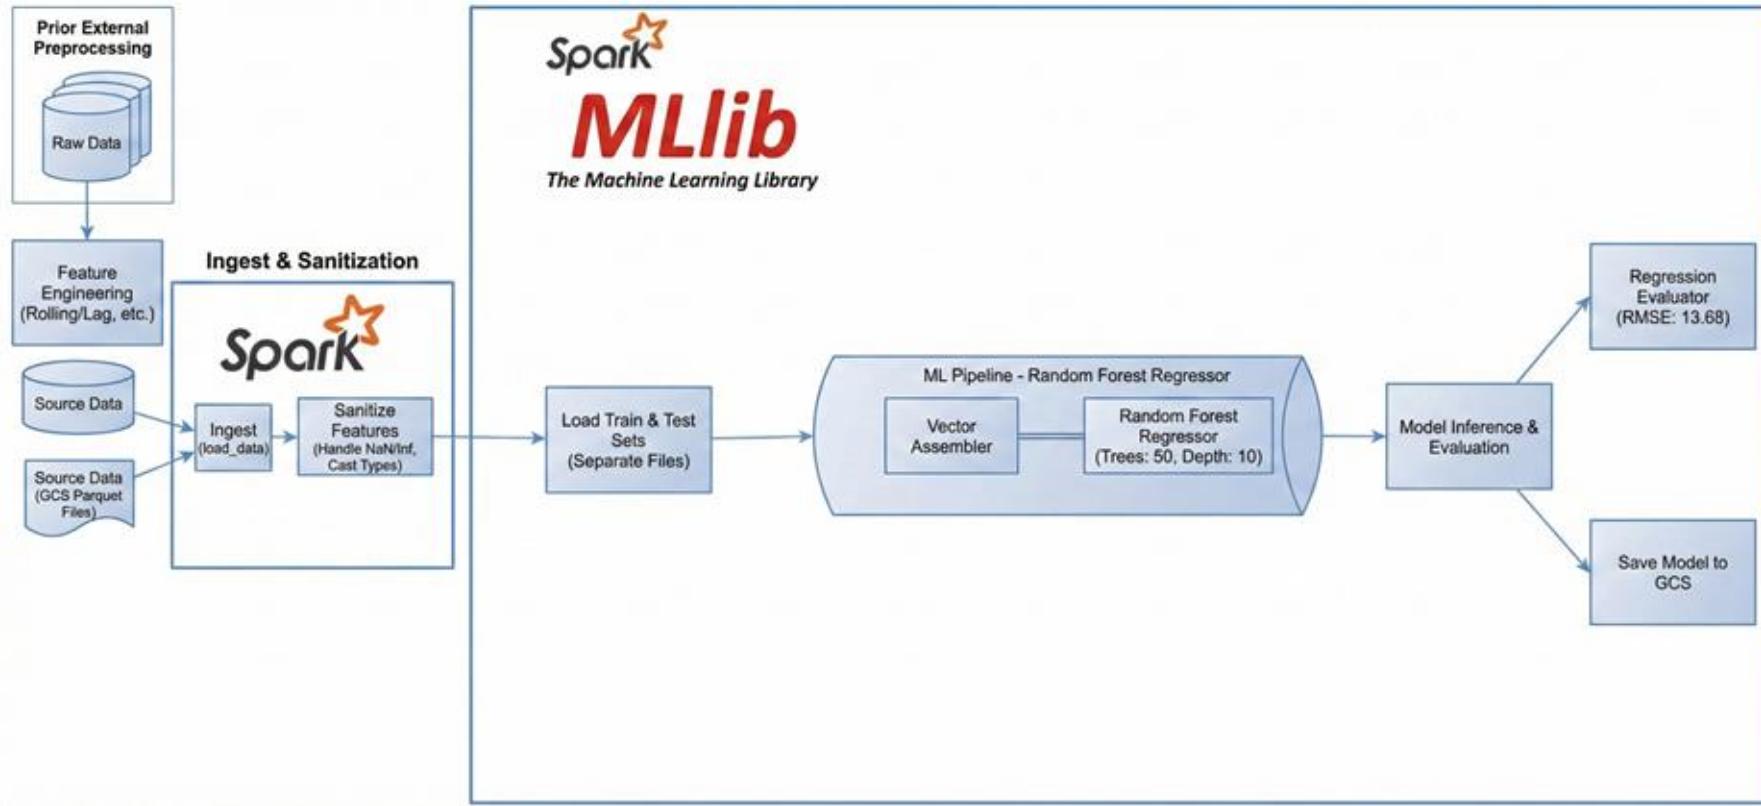
## 2. Total Cluster Capacity

- **Compute Power:** 16 vCPUs
- **Memory Pool:** 64 GB RAM

## 3. Memory Optimization

- **Why High-Memory Nodes?** Unlike Hadoop MapReduce (disk-based), Spark is an **In-Memory** engine.
- **Benefit:** Data remains cached in RAM between iterations, significantly accelerating iterative algorithms like **Random Forest**.





# Data Pipeline: Storage & Optimization

## 1. Storage Format: Parquet

- **Why Parquet:** Superior compression and strict schema preservation compared to CSV.
- **I/O Optimization:** Leverages Pushdown:
  - **Projection Pushdown:** Spark reads only required columns.
  - **Predicate Pushdown:** Filters are applied at the storage layer; irrelevant file blocks are skipped entirely.

## 2. Parallelism Strategy: Partitioning

- **Configuration:** `df.repartition(16)`
- **The Logic:** Cluster Capacity = 2 Workers × 8 Cores = **16 Total Cores**.
- **Optimization Goal:** Achieved a **1:1 Task-to-Core Ratio**.
  - Too Few (<16): Idle CPUs
  - Too Many (>16): High scheduler overhead

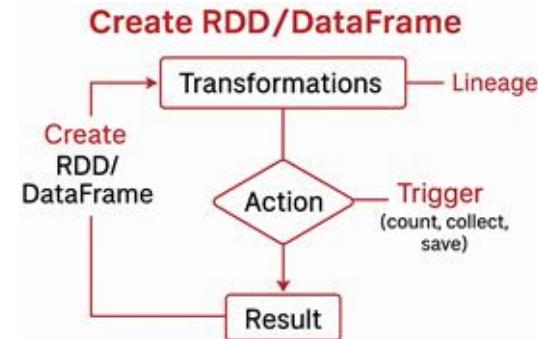


**Result: Maximum hardware utilization**

# Spark Execution & Model Preparation

## The Paradigm

- **Concept:** Transformations are not executed immediately. Spark records them as a Logical Plan.
- **Trigger:** Execution only occurs when an Action is called.
- **Benefit of Lazy Evaluation:** This delay allows Spark to optimize the plan before execution by reordering operations to minimize data movement and processing cost.



## Data Formatting

- **The Requirement:** Unlike Scikit-Learn which accepts Matrices or DataFrames, Spark MLlib requires features to be condensed into a Single Vector Column.
- **Implementation:** Used `VectorAssembler` to fuse individual feature columns (Store ID, Item ID, Sales Lags) into a compact feature vector for the model.

```
# 3. Build Pipeline
assembler = VectorAssembler(
    inputCols=feature_cols,
    outputCol="features",
    handleInvalid="keep"
)
```

*Transformation: [Col A, Col B, Col C] → [Vector(A, B, C)]*

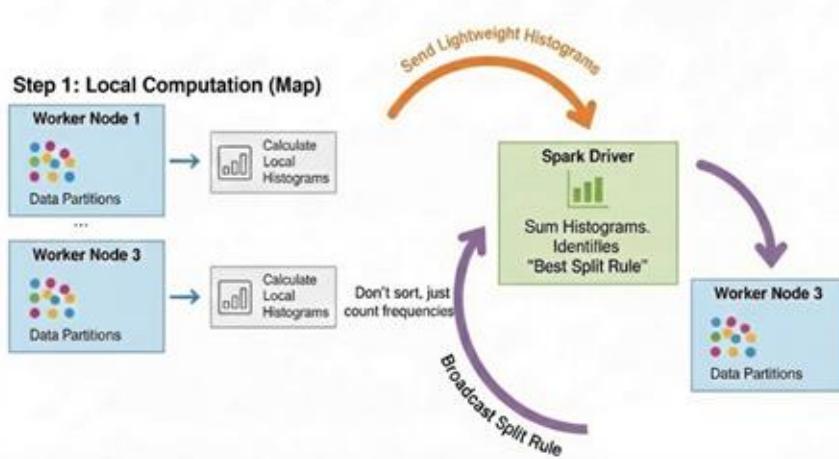
# Distributed Random Forest

Challenge	Constraint
 Single view of ALL data needed for optimal split.	 Data partitioned across worker nodes. Moving all data causes memory/network bottleneck

## Distributed Random Forest: The Binning

### The "Binning" Strategy

Uses Map-Reduce to approximate splits efficiently



### Problem:

Finding the **globally optimal split rule** for decision trees is **bottlenecked** by the need to **aggregate all data partitioned across worker nodes**.

### Solution:

- **Map:**

Worker nodes compute and send lightweight local histograms (compressed data summaries) instead of moving all data.

- **Reduce/Driver:**

The Spark Driver aggregates these histograms to find the best approximate split rule quickly.

- **Synchronization:**

The new split rule is broadcast back to all workers to continue tree growth.

### Result:

**Efficiently** approximates the optimal split with minimal network communication.

# Distributed Model Training

## 1. Stage 109 (Map): Parallel Computation

- **Data Parallelism:** Workers read partitions directly from disk.
- **Local Bootstrapping:** Independent sampling with replacement.

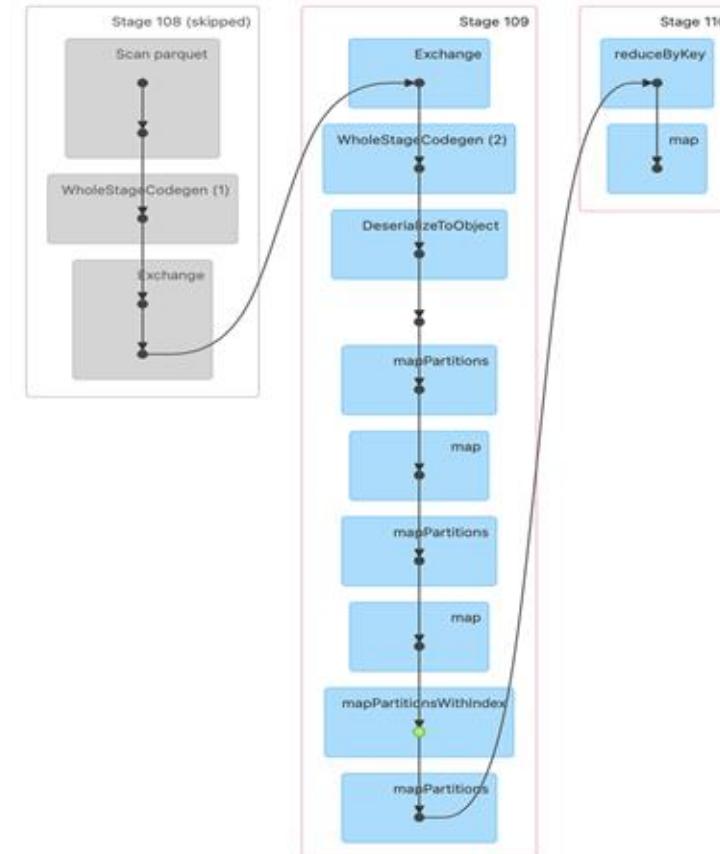
## 2. The Shuffle (Exchange)

- **Network Barrier:** Partial metadata transmitted to synchronize nodes.
- **Optimization:** Only summary statistics shuffled (no raw data).

## 3. Stage 110 (Reduce): Aggregation

- **Global Aggregation:** `reduceByKey` sums partial histograms.
- **Driver Sync:** `collectAsMap` returns the "best split" to the Driver to update the tree.

*Code moves to data, not the other way around. Heavy computation runs locally where data resides. Only lightweight summary statistics are shuffled, minimizing network latency.*



## Code Snippet - Main Logic (RF) - Training & Prediction

```
# 3. Build Pipeline
assembler = VectorAssembler(
    inputCols=feature_cols,
    outputCol="features",
    handleInvalid="keep"
)

rf = RandomForestRegressor(
    featuresCol="features",
    labelCol="unit_sales",
    numTrees=CONFIG["NUM TREES"],
    maxDepth=CONFIG["MAX DEPTH"],
    seed=CONFIG["SEED"]
)

pipeline = Pipeline(stages=[assembler, rf])

# 4. Train
print("Starting Random Forest training...") )
model = pipeline.fit(train_df)
print("Training completed.")

# 2. Clean Test Data (Must apply same cleaning as training!)
test_df, _ = sanitize_features(test_df, EXCLUDED_COLS)

# 3. Load Model
print(f"Loading model from {CONFIG['MODEL_OUTPUT_PATH']}...")
loaded_model = PipelineModel.load(CONFIG["MODEL_OUTPUT_PATH"])
print("Model loaded.")

# 4. Generate Predictions
print("Running inference...")
predictions = loaded_model.transform(test_df)

# 5. Calculate RMSE
evaluator = RegressionEvaluator(
    labelCol="unit_sales",
    predictionCol="prediction",
    metricName="rmse"
)
rmse = evaluator.evaluate(predictions)
print(f"Root Mean Squared Error (RMSE): {rmse:.4f}")
```

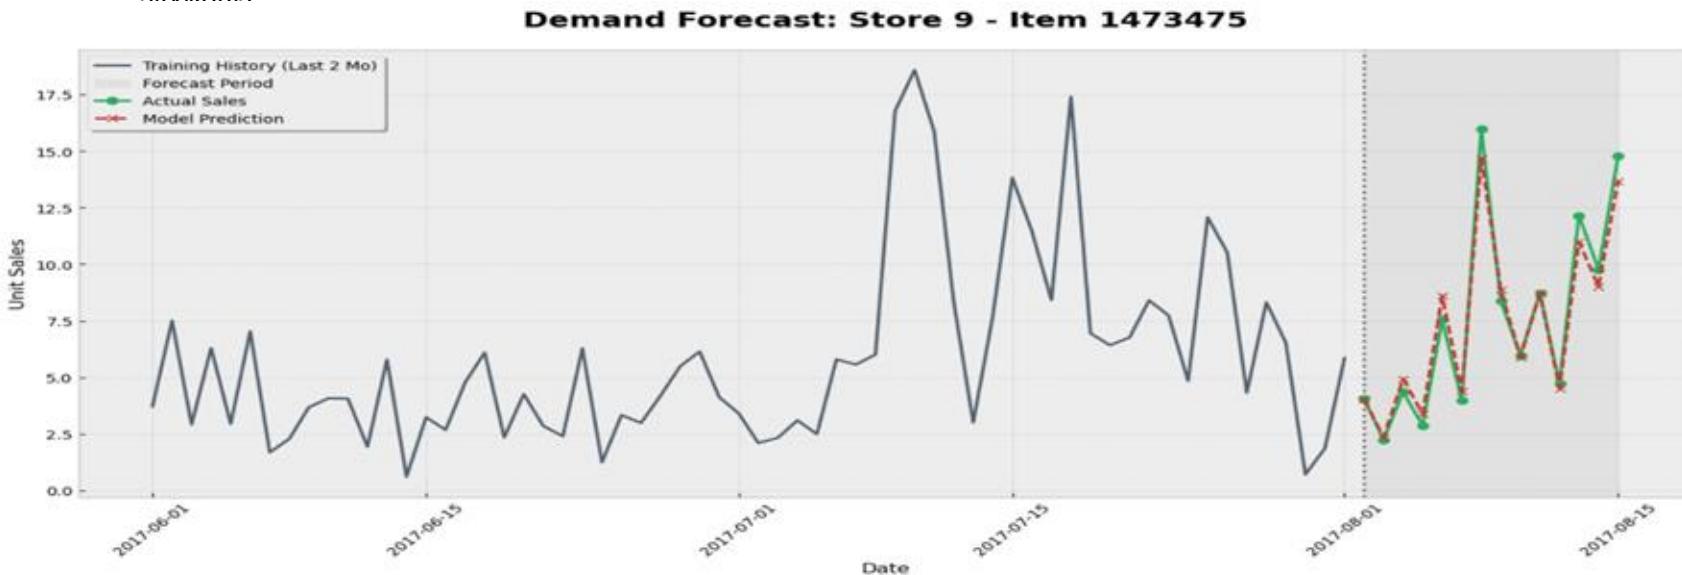
# Results & Discussion

## Model Performance

- Metric: RMSE  $\approx 13.69$ .
- Interpretation: On average, predictions deviate by  $\sim 13$  inventory units. This variance indicates the baseline for calculating minimum extra Stock requirements to prevent stockouts.

## Feature Importance

- Primary Driver: The Lagged Variables - `store_item_month_avg_sales`.
- Model Insight: The model identified historical sales behavior as the strongest predictor of future demand.



# THANK YOU